

# NEURAL NETWORK

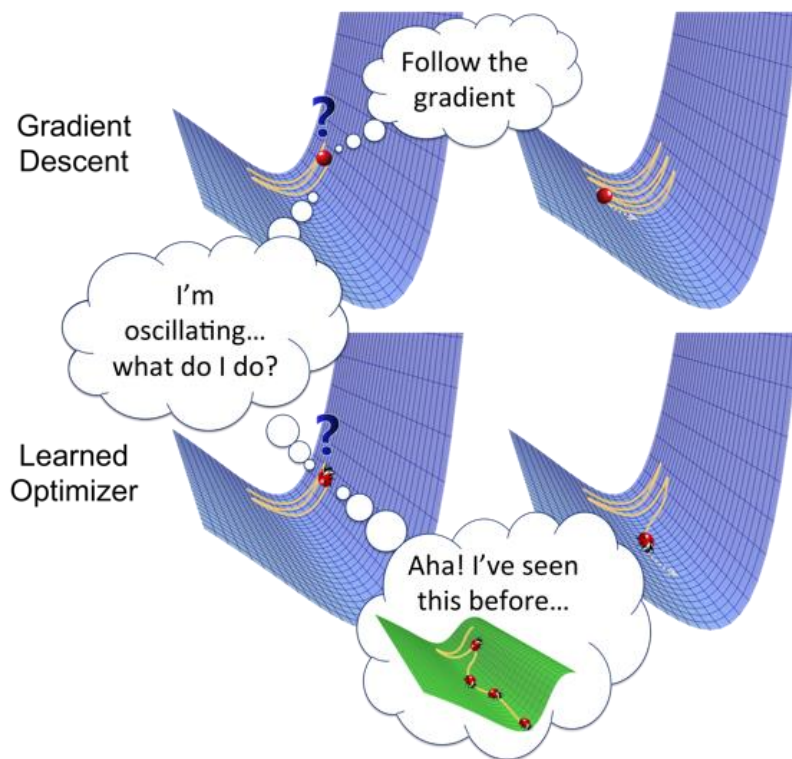
---

# Why Neural Networks now?


- Better activation functions for neural layers
- Better weight-initialization schemes
- Better optimization schemes
- Faster computations

# Optimizers and Parameters

- Optimization ~ change the weights and learning rate in order to reduce the loss function.
- More than one way ~ **Why do we need more?**



# Optimizers and Parameters




# Optimizers and Parameters

- Gradient Decent ~ The Most Basic
- Stochastic Gradient Descent
- Mini Batch SGD
- Momentum
  - GD:  $\theta = \theta - \alpha \cdot \nabla J(\theta)$
  - Mom:  $V(t) = \gamma V(t-1) + \alpha \cdot \nabla J(\theta)$ ,  $\theta = \theta - V(t)$ .


# Optimizers and Parameters

- Nesterov Accelerated Gradient




# Learning Rate

- Gradient Decent ~ Learning Rate
- Optimal value for the learning rate, eta ( $\eta$ ).
  - too small ~ too slow
  - too large ~ overshoot, no convergence
- No “optimal” solution



# Learning rate selection

The effects of step size (or “learning rate”)





# Learning Rate

- Fixed Learning Rate (usually small, say 0.1 or 0.01)
- Annealing Learning Rate
- Cyclical Learning Rate
- Adaptive Learning Rate

# Learning Rate

- How to determine?
- Start with a traditional default value
- Use Diagnostic plots
- Sensitivity Analysis (Hyperparameter tuning)
- Adaptive Learning Rate

# Momentum


- Including an exponentially weighted average of the prior updates while updating weights
- Motivation is to cause updating in the same direction in the future.
- Momentum is set to a value greater than 0.0 and less than one, where common values such as 0.5, 0.9 and 0.99 are used in practice.

# Optimizers and Parameters


- Adagrad
  - Non-constant learning rate
- AdaDelta
  - Improvement of Adagrad
  - Nondecaying learning rate

# Optimizers and Parameters


- RMSProp
- Adam
  - Adaptive Moment Estimation



# Optimizers and Parameters



# Optimizers and Parameters




# Batch Size Selection

- Batch size controls the accuracy of the estimate of the error gradient
- Batch, Stochastic, and Minibatch gradient descent
- Tradeoff in determining the batch size ~ the speed and stability of the learning process.



# Batch Size Selection




# Network Structure

- Number of layers and size of each layer in NN is important
- Used to have shallow networks but recently deep networks are popular
- Too large networks (both size and number of layers) tend to overfit, **why?**

# Weight Initialization

- Weight initialization ~ updating the parameters normally
- Vanishing or exploding gradients



# Weight Initialization


- Zero Initialization
- Random Initialization
- Xavier Initialization
- He Initialization

# Preprocessing the Data


- Scaling the input data is important
- Large Values ~ Large Weights ~ Unstable model
- Rule of Thumb ~ range 0-1 or standardized with zero mean and stan. dev. of 1.

# Scaling the Data


Unnormalized



Normalized



# Preprocessing the Data



(Assume  $X$  [NxD] is data matrix,  
each example in a row)

# Preprocessing the Data

- Data Normalization



# Preprocessing the Data


- Data Standardization

# Overfitting

- Deep neural networks require lots of data, and can overfit easily
- The more weights you need to learn, the more data you need
- That's why with a deeper network, you need more data for training than for a shallower network
- Ways to prevent overfitting include:
  - Using a validation set to stop training or pick parameters
  - Regularization
  - Data Augmentation

# Over-training prevention


- Running too many epochs can result in over-fitting.



- Keep a hold-out validation set and test accuracy on it after every epoch. Stop training when additional epochs actually increase validation error.


# Determining best number of hidden units

- Too few hidden units prevents the network from adequately fitting the data.
- Too many hidden units can result in over-fitting.



- Use internal cross-validation to empirically determine an optimal number of hidden units.

# Effect of number of neurons



↑  
more neurons = more capacity


# Effect of regularization

Do not use size of neural network as a regularizer. Use stronger regularization instead:


$\lambda = 0.001$



$\lambda = 0.01$



$\lambda = 0.1$



(you can play with this demo over at ConvNetJS: <http://cs.stanford.edu/people/karpathy/convnetjs/demo/classify2d.html>)

# Weight Regularization $\lambda$ = regularization strength (hyperparameter)

$$L(W) = \underbrace{\frac{1}{N} \sum_{i=1}^N L_i(f(x_i, W), y_i)}_{\text{Data loss}} + \underbrace{\lambda R(W)}_{\text{Regularization}}$$

**Data loss:** Model predictions should match training data

**Regularization:** Prevent the model from doing *too* well on training data

## Simple examples

L2 regularization:  $R(W) = \sum_k \sum_l W_{k,l}^2$

L1 regularization:  $R(W) = \sum_k \sum_l |W_{k,l}|$

Elastic net (L1 + L2):  $R(W) = \sum_k \sum_l \beta W_{k,l}^2 + |W_{k,l}|$

## More complex:

Dropout


Batch normalization

Why regularize?

- Express preferences over weights
- Make the model *simple* so it works on test data

# Weight Regularization

- Preferring simple models




Regularization pushes against fitting the data *too* well so we don't fit noise in the data




# Dropout


- Dropout
  - Randomly turn off some neurons
  - Allows individual neurons to independently be responsible for performance



(a) Standard Neural Net



(b) After applying dropout.



Dropout: A simple way to prevent neural networks from overfitting

# Batch Normalization

**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_{1\dots m}\}$ ;

Parameters to be learned:  $\gamma, \beta$

**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$


$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

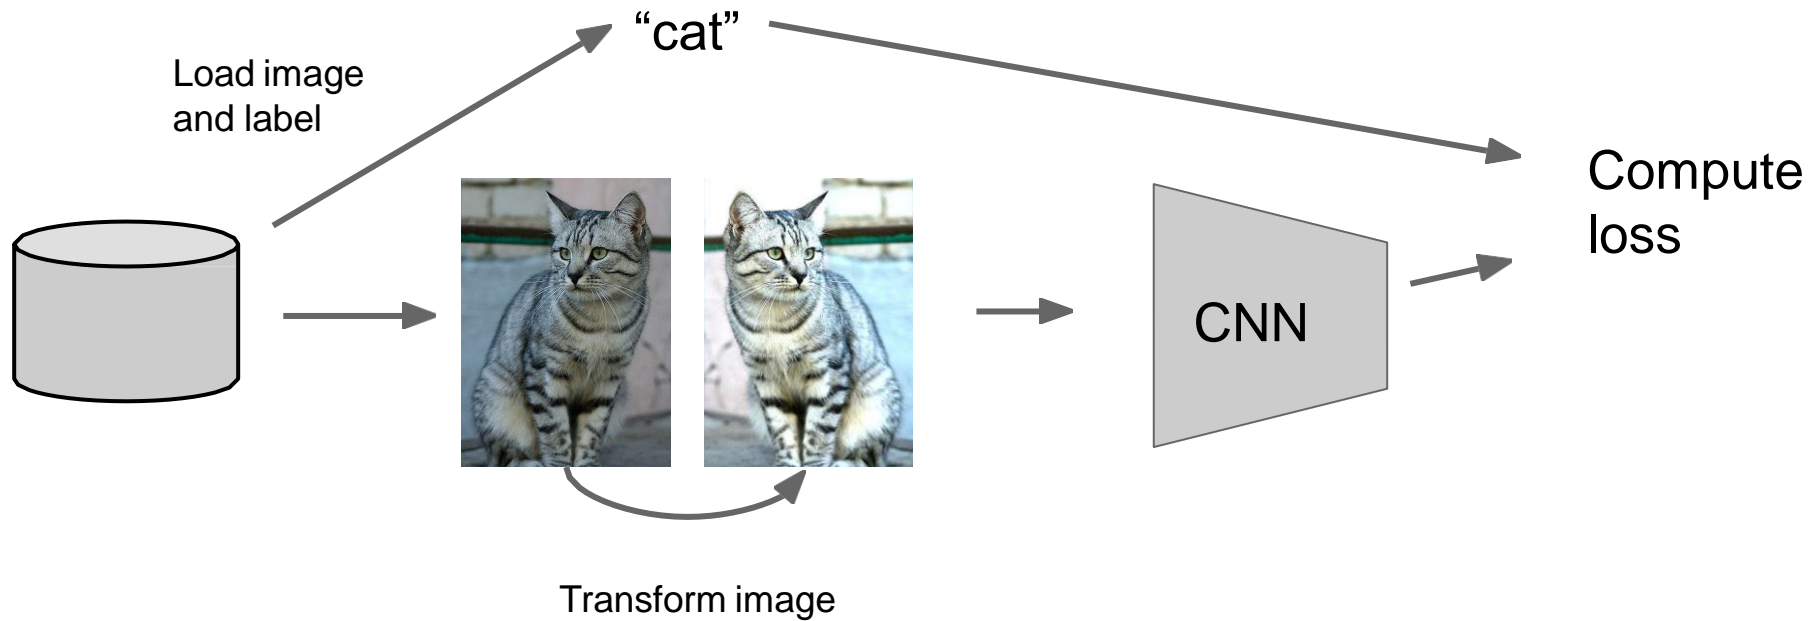
$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

- Improves gradient flow through the network
- Allows higher learning rates
- Reduces the strong dependence on initialization
- Acts as a form of regularization

# Data Augmentation



# Data Augmentation




# Data Augmentation

- Horizontal Flips

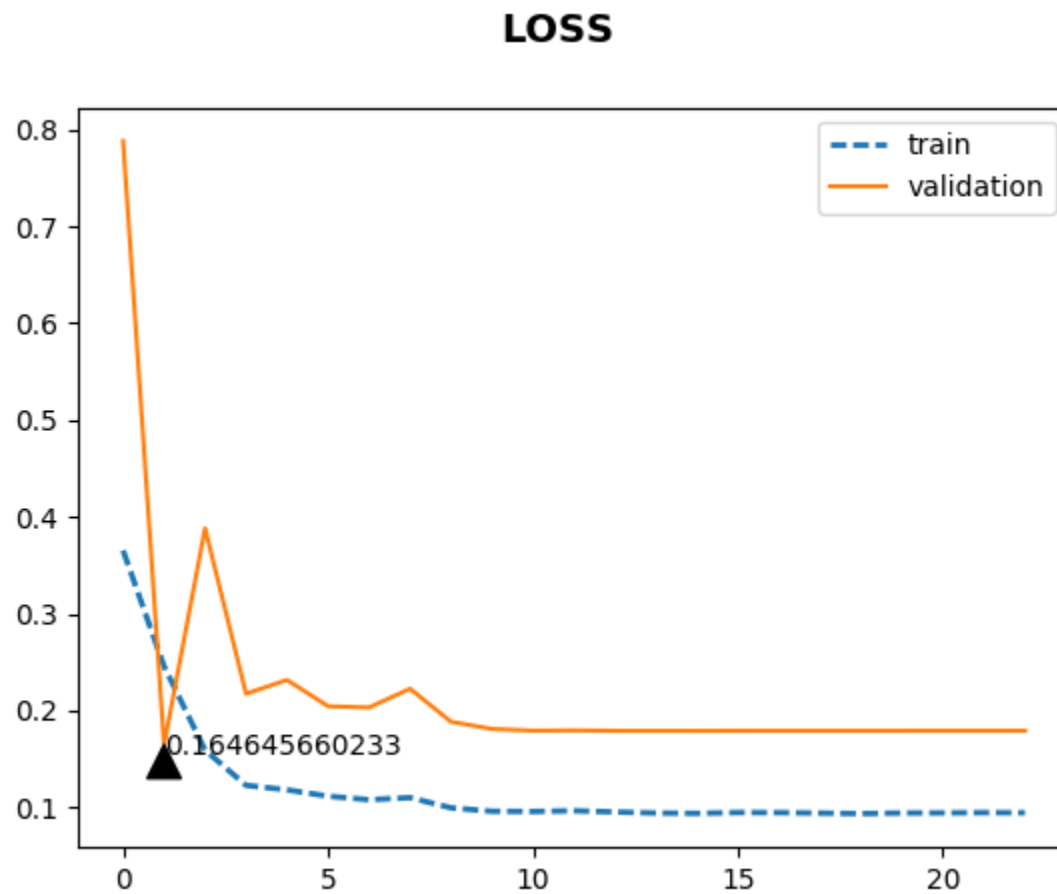


# Data Augmentation

- Get creative for your problem!
- Random mix/combinations of :
  - translation
  - rotation
  - stretching
  - shearing,
  - lens distortions
  - ...



# Callbacks



# CPU vs GPU

	Cores	Clock Speed	Memory	Price	Speed
<b>CPU</b> (Intel Core i7-7700k)	4 (8 threads with hyperthreading)	4.2 GHz	System RAM	\$385	~540 GFLOPs FP32
<b>GPU</b> (NVIDIA RTX 2080 Ti)	3584	1.6 GHz	11 GB GDDR6	\$1199	~13.4 TFLOPs FP32
<b>TPU</b> NVIDIA TITAN V	5120 CUDA, 640 Tensor	1.5 GHz	12GB HBM2	\$2999	~14 TFLOPs FP32 ~112 TFLOP FP16
<b>TPU</b> Google Cloud TPU	?	?	64 GB HBM	\$4.50 per hour	~180 TFLOP

**CPU:** Fewer cores, but each core is much faster and much more capable; great at sequential tasks


**GPU:** More cores, but each core is much slower and “dumber”; great for parallel tasks

**TPU:** Specialized hardware for deep learning



# CPU vs GPU in practice


(CPU performance not well-optimized, a little unfair)



Data from <https://github.com/jcjohnson/cnn-benchmarks>

# CPU / GPU Communication

Model  
is here



Data is here

If you aren't careful, training can bottleneck on reading data and transferring to GPU!

## Solutions:

- Read all data into RAM
- Use SSD instead of HDD
- Use multiple CPU threads to prefetch data

# Training: Best practices

- Center (subtract mean from) your data
- To initialize weights, use “Xavier or He initialization”
- Use RELU or leaky RELU or ELU, don’t use sigmoid
- Use mini-batch
- Use data augmentation
- Use regularization
- Use batch normalization
- Use cross-validation for your parameters
- Learning rate: too high? Too low?